

Project 1

Wiki Documentation

This project focuses on the implementation of a system scheduler in a RISC-V xv6 architecture. Given that a Round-Robin (RR) scheduler is used by default, the objective of this project is to implement a First-Come, First-Served (FCFS) and Multi-Level Feedback Queue (MLFQ) scheduler. In this case, there are five system calls to be implemented, with each serving a purpose in implementing the two schedulers.

Given that the built-in scheduler function in the `proc.c` file uses a RR scheduler, I decided to reference its implementation merely as a guide in defining and implementing the FCFS and MLFQ schedulers. I approached this assignment by implementing the following system calls: `yield()`, `getlev()`, `setpriority()`, `fcfsmode()`, and `mlfqmode()` by modifying the following files: `sysproc.c`, `syscall.c`, `syscall.h`, `defs.h`, `user.h`, `user.pl`, `proc.c`, and `proc.h`. Below are the following steps I took to implement these.

Step 1: Define the system calls in the `sysproc.c` file

In order to create a new system call, it should be defined in the `sysproc.c` file as the file specifies the implementation of these system calls and overall contains the system calls that handle processes and its management.

I implemented five system calls: `sys_yield(void)`, `sys_getlev(void)`, `sys_setpriority(void)`, `sys_fcfsmode(void)`, and `sys_mlfqmode(void)`. Each serves a different role in implementing system schedulers. Since the `yield()` function exists as a built-in function, I called the function in the `syscall`. For `getlev` and `set priority`, I made it so that `getlev` returns the queue level of the process and `setpriority` sets the priority of the process with the `pid`. Additionally, the `argint()` call in `sys_setpriority()` extracts integer arguments passed from a user program to the kernel `syscall`, thus it needed to be changed to return a value (`void` \rightarrow `int`). More specific implementations of the functions were made in the `proc.c` file.

```
9  extern int current_mode;
10 uint64
11 sys_yield(void)
12 {
13     yield();
14     return 0;
15 }
16
17 uint64
18 sys_getlev(void)
19 {
20     struct proc *p = myproc();
21     if (current_mode == FCFS)
22         return 99;
23     return p->queue_level;
24 }
25
26 uint64
27 sys_setpriority(void)
28 {
29     int pid, priority;
30     if(argint(0, &pid) < 0 || argint(1, &priority) < 0)
31         return -1;
32     return setpriority(pid, priority);
33 }
34
35
36 uint64
37 sys_mlfqmode(void)
38 {
39     return mlfqmode();
40 }
41
42 uint64
43 sys_fcfsmode(void)
44 {
45     return fcfsmode();
46 }
```

Figure 1. `sysproc.c` file

Step 2: Define the functions in the defs.h file

I defined the functions in the defs.h file so that I could access these files when initializing in the syscall.c file. Since 'yield' is an in-built function, I just created four definitions for the remaining four functions. Also, I changed the argint(int, int*) from void to an int function so that it returned a value when used in the setpriority() system call.

```
105 void        yield(void);
106 int         either_copyout(int user_dst, uint64 dst, void *src, uint64 len);
107 int         either_copyin(void *dst, int user_src, uint64 src, uint64 len);
108 void        procdump(void);
109 int         getlev(void);
110 int         setpriority(int pid, int priority);
111 int         mlfqmode(void);
112 int         fcfsmode(void);
113
```

```
// syscall.c
int         argint(int, int*);
```

Figure 2. defs.h file

Step 3: Register the system calls in the syscall.c file

I made a few changes to the syscall.c file, which added/registered the new functions as new system calls. Moreover, I changed the argint() function to an int function and included a "return 0" at the end so that it can be used when it is called in the setpriority function.

```
105 extern uint64 sys_yield(void);          [SYS_yield]    sys_yield,
106 extern uint64 sys_getlev(void);         [SYS_getlev]   sys_getlev,
107 extern uint64 sys_setpriority(void);     [SYS_setpriority] sys_setpriority,
108 extern uint64 sys_mlfqmode(void);        [SYS_mlfqmode] sys_mlfqmode,
109 extern uint64 sys_fcfsmode(void);       [SYS_fcfsmode] sys_fcfsmode,
```

```
55 // Fetch the nth 32-bit system call argument.
56 int
57 argint(int n, int *ip)
58 {
59     *ip = argraw(n);
60     return 0;
61 }
```

Figure 3. syscall.c file

Step 4: Register the functions in the system call numbers

In addition to registering the functions in the syscall.c file, I also registered it in the syscall.h file by defining its system call number to the latest available number, which in this case is 22-26. This is done so that a number corresponds to each specific system call.

```
23 #define SYS_yield 22
24 #define SYS_getlev 23
25 #define SYS_setpriority 24
26 #define SYS_mlfqmode 25
27 #define SYS_fcfsmode 26
```

Figure 4. syscall.h file

Step 5: Initializing variables in the proc.h file

In order to use certain variables in the proc.c file, I initialized the following variables in the proc.h file. I defined RR as 0, FCFS as 1, and MLFQ as 2, and I defined an extern variable `current_mode` that indicates which scheduler mode is being used.

```
28 #define RR 0
29 #define FCFS 1
30 #define MLFQ 2
31 extern struct cpu cpus[NCPU];
32 extern int current_mode;
96 int priority;
97 int queue_level;
98 int ctime;
99 int time_quantum;
100 int ticks;
```

Figure 5. proc.h file

Step 6: Adding required specifications to the functions in the proc.c file

Several changes were made to the proc.c file. First of all, I initialized the `current_mode` as FCFS in order to schedule processes within FCFS after the first xv6 boot. I also created a `boost_ticks` variable that will be used in the MLFQ scheduler.

```
16 int current_mode = FCFS;
17 struct spinlock pid_lock;
18 int boost_ticks = 0;
```

Figure 6. proc.c file (initializing variables)

Next, I added additional initializations to the `allocproc` function that will be used in the scheduler. Here, `ctime` helps manage process scheduling based on their creation time, `queue_level` and `priority` allow for adjustments based on the process's behavior, and `time_quantum` is used to manage how long processes run before being preempted.

```
145 // Set up new context to start executing at forkret,
146 // which returns to user space.
147 memset(&p->context, 0, sizeof(p->context));
148 p->context.ra = (uint64)forkret;
149 p->context.sp = p->kstack + PGSIZE;
150 p->ctime = ticks;
151 p->queue_level = 0;
152 p->priority = 3;
153 p->time_quantum = 0;
154 return p;
155 }
```

Figure 7. proc.c file (allocproc() function)

Afterwards, I implemented the following functions: `setpriority`, `mlfqmode`, and `fcfsmode`. `setpriority()` returns 0 when priority of the process was set with the given pid properly, -1 if there's no process with the given pid, and -2 if the priority value is not between 0 and 3. `mlfqmode()` and `fcfsmode()` returns 0 when it successfully transitions modes, and -1 if the system is already in MLFQ/FCFS mode.

```

789  int
790  setpriority(int pid, int priority)
791  {
792      struct proc *p;
793      if (priority <= 0 || priority >= 3) {
794          return -2;
795      }
796      for (p = proc; p < &proc[NPROC]; p++) {
797          acquire(&p->lock);
798          if (p->pid == pid) {
799              p->priority = priority;
800              release(&p->lock);
801              return 0;
802          }
803          release(&p->lock);
804      }
805      return -1;
806  }

897  int mlfqmode(void) {
898      if (current_mode == MLFQ) {
899          return -1;
900      }
901      current_mode = MLFQ; //switch
902      ticks = 0; //reset ticks
903      for (struct proc *p = proc; p < &proc[NPROC]; p++) {
904          acquire(&p->lock);
905          if (p->state != UNUSED) {
906              p->queue_level = 0; //Lo queue
907              p->priority = 3; //reset
908              p->time_quantum = 0; //reset
909          }
910          release(&p->lock);
911      }
912      return 0;
913  }

915  int fcfsmode(void) {
916      if (current_mode == FCFS) {
917          return -1;
918      }
919      current_mode = FCFS;
920      ticks = 0;
921      for (struct proc *p = proc; p < &proc[NPROC]; p++) {
922          acquire(&p->lock);
923          if (p->state != UNUSED) {
924              p->queue_level = 0;
925              p->priority = 3;
926              p->time_quantum = 0;
927          }
928          release(&p->lock);
929      }
930      return 0;
931  }

```

Figure 8. *proc.c* file (`setpriority()`, `fcfsmode()`, `mlfqmode()` function)

The scheduler() function runs in an infinite loop (for(;;)) to continuously check and schedule processes. The intr_on() enables interrupts to prevent deadlocks. It switches between three different scheduler modes: RR, FCFS, and MLFQ. It first checks if the current_mode is the RR mode, which is the default scheduler of xv6. It then checks each process to see if it is runnable and if a process is found, it is switched to the RUNNING state and the context is switched to that process. If no runnable process is found, the CPU waits for an interrupt using asm volatile("wfi").

```
449 void
450 scheduler(void)
451 {
452     struct proc *p;
453     struct cpu *c = mycpu();
454
455     c->proc = 0;
456     for(;;){
457         intr_on();
458         if (current_mode == RR){
459             // The most recent process to run may have had interrupts
460             // turned off; enable them to avoid a deadlock if all
461             // processes are waiting.
462             int found = 0;
463             for(p = proc; p < &proc[NPROC]; p++) {
464                 acquire(&p->lock);
465                 if(p->state == RUNNABLE) {
466                     // Switch to chosen process. It is the process's job
467                     // to release its lock and then reacquire it
468                     // before jumping back to us.
469                     p->state = RUNNING;
470                     c->proc = p;
471                     swtch(&c->context, &p->context);
472
473                     // Process is done running for now.
474                     // It should have changed its p->state before coming back.
475                     c->proc = 0;
476                     found = 1;
477                 }
478                 release(&p->lock);
479             }
480             if(found == 0) {
481                 // nothing to run; stop running on this core until an interrupt.
482                 intr_on();
483                 asm volatile("wfi");
484             }
485         }
486     }
487 }
```

Figure 9. proc.c file (scheduler() RR)

Next, it checks if the `current_mode` is the FCFS mode. Processes are executed in the order they are created based on their creation time (`ctime`). The loop iterates over the processes and selects the one with the minimum creation time, which is the process created the earliest. If no process is runnable, the CPU waits for an interrupt.

```
486     else if (current_mode == FCFS) {
487         struct proc *minp = 0; //minimum process
488         for (p = proc; p < &proc[NPROC]; p++) {
489             acquire(&p->lock);
490             if (p->state == RUNNABLE) {
491                 if (minp == 0 || p->ctime < minp->ctime) {
492                     if (minp) {
493                         release(&minp->lock);
494                     }
495                     minp = p;
496                 }
497                 else {
498                     release(&p->lock);
499                 }
500             }
501             else {
502                 release(&p->lock);
503             }
504         }
505         if (minp) {
506             minp->state = RUNNING;
507             c->proc = minp;
508             swtch(&c->context, &minp->context);
509             c->proc = 0;
510             release(&minp->lock);
511         }
512         else {
513             intr_on();
514             asm volatile("wfi");
515         }
516     }
```

Figure 10. `proc.c` file (`scheduler()` FCFS)

Lastly, it checks if the `current_mode` is the MLFQ mode. It uses multiple queues, each representing a different priority level. Processes are initially placed in higher-priority queues and are demoted to lower-priority queues if they exhaust their time slices. In order to prevent starvation, all processes are boosted by resetting their queue level to 0 (highest priority) every 50 ticks (priority boosting). The scheduler looks for runnable processes in the highest priority queue (level 0) and if no processes are found in level 0, it checks level 1 then level 2. If a process uses up its time slice, it is demoted to a lower priority level and if no processes are found in any of the levels, the system defaults to FCFS by calling `fcfsmode()`.

```

518     else if (current_mode == MLFQ) {
519         struct proc *p;
520         struct proc *chosen = 0;
521         int level;
522
523         boost_ticks++;
524         if (boost_ticks >= 50) { //priority boost every 50 ticks
525             for(p = proc; p < &proc[NPROC]; p++) { //reset after 50 ticks
526                 acquire(&p->lock);
527                 if (p->state == RUNNABLE || p->state == RUNNING) {
528                     p->queue_level = 0;
529                     p->ticks = 0;
530                 }
531                 release(&p->lock);
532             }
533             boost_ticks = 0;
534         }
535         for (level = 0; level <= 2; level++) { //checking each queue level
536             for(p = proc; p < &proc[NPROC]; p++) {
537                 acquire(&p->lock);
538                 if (p->state == RUNNABLE && p->queue_level == level) {
539                     chosen = p;
540                     goto found_mlfq; //if process found -> exit loop
541                 }
542                 release(&p->lock);
543             }
544         }
545         fcfsmode();
546
547         found_mlfq:
548         if (chosen) {
549             chosen->state = RUNNING;
550             c->proc = chosen;
551             chosen->ticks++; //increment ticks
552
553             switch(&c->context, &chosen->context); //context switch
554             c->proc = 0;
555
556             if (chosen->state == RUNNABLE) { //demote to lower level if exhausted time slice
557                 if ((chosen->queue_level == 0 && chosen->ticks >= 1) ||
558                     (chosen->queue_level == 1 && chosen->ticks >= 2) ||
559                     (chosen->queue_level == 2 && chosen->ticks >= 4)) {
560                     if (chosen->queue_level < 2) {
561                         chosen->queue_level++;
562                     }
563                     chosen->ticks = 0; //reset
564                 }
565             }
566             release(&chosen->lock);
567         }
568         else {
569             intr_on();
570             asm volatile("wfi");
571         }
572     }
573 }
574 }

```

Figure 11. `proc.c` file (`scheduler()` MLFQ)

Step 7: Declare the function in the user.h file

In order to make the system call available at the user program, I declared the function in the user/user.h file.

```
25 void yield(void);
26 int getlev(void);
27 int setpriority(int, int);
28 int mlfqmode(void);
29 int fcfsmode(void);
```

Figure 12. user.h file

Step 8: Add a new macro in the usys.pl file

Afterward, I added new macros in the usys.pl file so that it generates a system call stub that allows user programs to invoke these system calls. The usys.pl file overall defines how to preprocess when using the system call in a user program.

```
39 entry("yield");
40 entry("getlev");
41 entry("setpriority");
42 entry("mlfqmode");
43 entry("fcfsmode");
```

Figure 13. usys.pl file

Step 9: Add the code file to the makefile

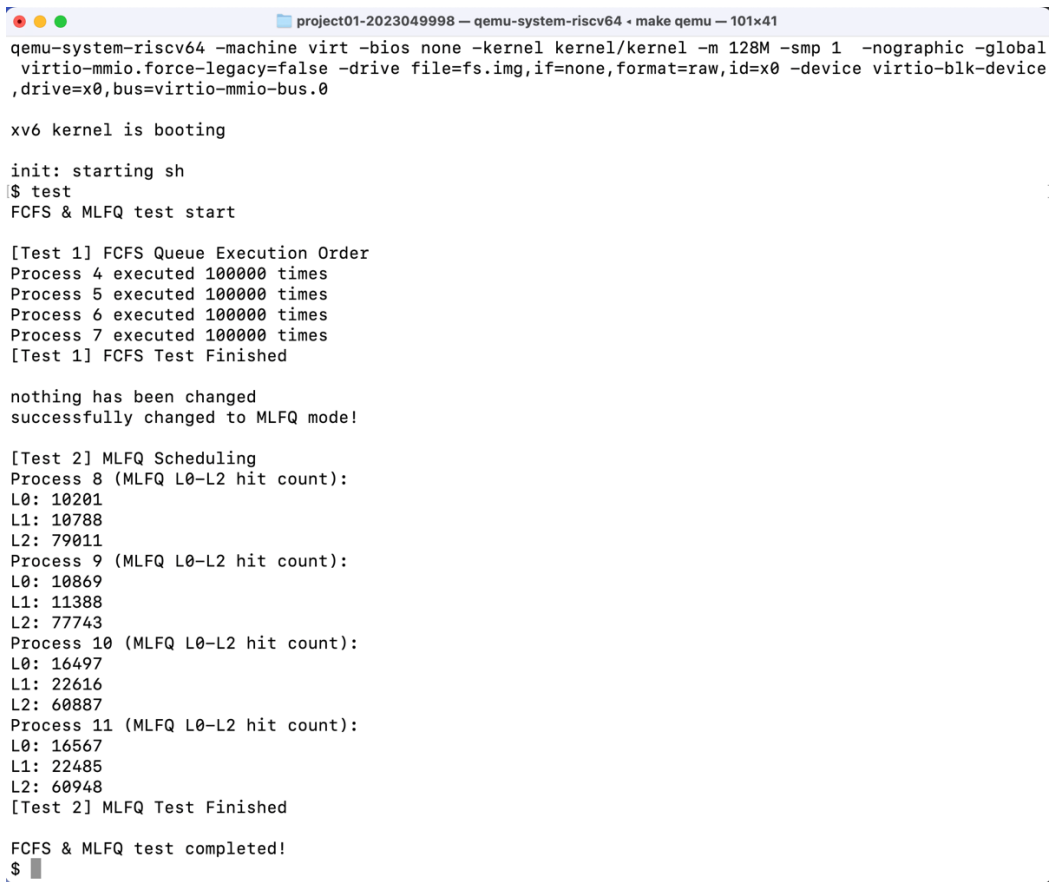
Adding the code `$U/_test\` into the makefile ensures that the user program test.c is properly compiled and linked once I run the makefile.

```
125 UPROGS=\
126     $U/_cat\
127     $U/_echo\
128     $U/_forktest\
129     $U/_grep\
130     $U/_init\
131     $U/_kill\
132     $U/_ln\
133     $U/_ls\
134     $U/_mkdir\
135     $U/_rm\
136     $U/_sh\
137     $U/_stressfs\
138     $U/_usertests\
139     $U/_grind\
140     $U/_wc\
141     $U/_zombie\
142     $U/_test\
```

Figure 14. Implementing the function

RESULTS

After implementing the `getppid()` function as a system call and `ppid.c` user program, I booted the xv6 kernel to execute the `ppid` user program. The results are as follows:



```
qemu-system-riscv64 -machine virt -bios none -kernel kernel/kernel -m 128M -smp 1 -nographic -global virtio-mmio.force-legacy=false -drive file=fs.img,if=none,format=raw,id=x0 -device virtio-blk-device,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

init: starting sh
$ test
FCFS & MLFQ test start

[Test 1] FCFS Queue Execution Order
Process 4 executed 100000 times
Process 5 executed 100000 times
Process 6 executed 100000 times
Process 7 executed 100000 times
[Test 1] FCFS Test Finished

nothing has been changed
successfully changed to MLFQ mode!

[Test 2] MLFQ Scheduling
Process 8 (MLFQ L0-L2 hit count):
L0: 10201
L1: 10788
L2: 79011
Process 9 (MLFQ L0-L2 hit count):
L0: 10869
L1: 11388
L2: 77743
Process 10 (MLFQ L0-L2 hit count):
L0: 16497
L1: 22616
L2: 60887
Process 11 (MLFQ L0-L2 hit count):
L0: 16567
L1: 22485
L2: 60948
[Test 2] MLFQ Test Finished

FCFS & MLFQ test completed!
$
```

Figure 15. Execution results

As seen in the photo, the program tests the FCFS scheduling, switches modes to MLFQ, then testing the MLFQ scheduling.

As for the flow of operation, given the `test.c` file, it starts with the `fork()` process being called to create four child processes that involves incrementing its own counter until it reaches `NUM_LOOP` (= 100,000). Here, the FCFS makes it so that each process is run until its complete before starting the next one, and it also means that processes are run in their PID/creation order.

Afterwards, the mode checks to see if it is in FCFS mode then if it is in MLFQ mode. Since it was in FCFS mode to begin with (because `current_mode` is assigned to be FCFS when the xv6 is booted), it prints "nothing has been changed". Then, it switches to MLFQ mode and prints out the confirmation.

In the MLFQ scheduler, it creates four child process then it calls `getlev()` to check its current level. It increments `count[level]` based on the result then it prints how many times it was scheduled at each level.

TROUBLESHOOTING

When trying to run the test.c file twice in a row, I was getting a message that despite starting and testing the FCFS mode, it says “successfully changed to FCFS mode!” before switching to MLFQ mode, which shouldn’t be the case considering that after MLFQ is tested and a new test is called, it should switch back to FCFS mode.

```
project01-2023049998 — qemu-system-riscv64 - make qemu — 101x34
$ test
FCFS & MLFQ test start

[Test 1] FCFS Queue Execution Order
Process 13 executed 100000 times
Process 14 executed 100000 times
Process 15 executed 100000 times
Process 16 executed 100000 times
[Test 1] FCFS Test Finished

successfully changed to FCFS mode!
successfully changed to MLFQ mode!

[Test 2] MLFQ Scheduling
Process 17 (MLFQ L0-L2 hit count):
L0: 1054
L1: 11401
L2: 87543
Process 18 (MLFQ L0-L2 hit count):
L0: 5635
L1: 11368
L2: 82997
Process 19 (MLFQ L0-L2 hit count):
L0: 11408
L1: 22483
L2: 66189
Process 20 (MLFQ L0-L2 hit count):
L0: 11415
L1: 22536
L2: 60849
[Test 2] MLFQ Test Finished

FCFS & MLFQ test completed!
$
```

I tried to fix this issue by calling `fcfsmode()` after `mlfq` scheduler is run, given the specification that “the `fcfsmode` system call is invoked in MLFQ mode, processes will be scheduled within the FCFS again.”

```
534         for (level = 0; level <= 2; level++) { //checking eqch queue level
535             for(p = proc; p < &proc[INPROC]; p++) {
536                 acquire(&p->lock);
537                 if (p->state == RUNNABLE && p->queue_level == level) {
538                     chosen = p;
539                     goto found_mlfq; //if process found -> exit loop
540                 }
541                 release(&p->lock);
542             }
543         }
544         fcfsmode();
545     }
```

By doing so, when running the test consecutively, it prints out “nothing has been changed” before switching to MLFQ mode.

<pre>xv6 kernel is booting init: starting sh \$ test FCFS & MLFQ test start [Test 1] FCFS Queue Execution Order Process 4 executed 100000 times Process 5 executed 100000 times Process 6 executed 100000 times Process 7 executed 100000 times [Test 1] FCFS Test Finished nothing has been changed successfully changed to MLFQ mode! [Test 2] MLFQ Scheduling Process 8 (MLFQ L0-L2 hit count): L0: 1054 L1: 10546 L2: 88400 Process 9 (MLFQ L0-L2 hit count): L0: 5086 L1: 10617 L2: 84297 Process 10 (MLFQ L0-L2 hit count): L0: 7988 L1: 20562 L2: 71450 Process 11 (MLFQ L0-L2 hit count): L0: 10578 L1: 20682 L2: 68740 [Test 2] MLFQ Test Finished FCFS & MLFQ test completed!</pre>	<pre>\$ test FCFS & MLFQ test start [Test 1] FCFS Queue Execution Order Process 13 executed 100000 times Process 14 executed 100000 times Process 15 executed 100000 times Process 16 executed 100000 times [Test 1] FCFS Test Finished nothing has been changed successfully changed to MLFQ mode! [Test 2] MLFQ Scheduling Process 17 (MLFQ L0-L2 hit count): L0: 9405 L1: 21295 L2: 69300 Process 18 (MLFQ L0-L2 hit count): L0: 10505 L1: 21222 L2: 68273 Process 19 (MLFQ L0-L2 hit count): L0: 15927 L1: 32015 L2: 52058 Process 20 (MLFQ L0-L2 hit count): L0: 15751 L1: 26301 L2: 57948 [Test 2] MLFQ Test Finished FCFS & MLFQ test completed!</pre>
--	---

Figure 16. test 1 (left) and test 2 (right)